

Class 6: Decision-Making and Interfaces for ML Models

Master Course:

Data-driven Systems Engineering (ML Operations)

440MI and 305SM

Agenda


- Learning Goals
- From Models to Decisions
- Visualizing Model Predictions
- Interfaces for Machine Learning
- Introduction to Streamlit
- Live Demo: Building a Streamlit
- Discussion and Best Practices

Learning Goals

- Understand how predictions from ML models can support decision-making.
- Visualize model outputs interactively.
- Build simple web interfaces for ML using Streamlit.
- Translate technical predictions into actionable insights.

MACHINE LEARNING ENGINEERING


DATA PIPELINE



EXPLORATION
&
VALIDATION

- Profiling
- "JUnit4Data"

WRANGLING
(CLEANING)



- Data versioning

TRAIN

TEST

MODEL
ENGINEERING

- Feature engineering
- Hyperparameters tuning

MODEL
EVALUATION

- Best model selection
 - Model performance metrics
 - accuracy
 - precision
 - recall
- F₁

MODEL
PACKAGING

- Model format
- ONNX
- JAR
- .pkl

MODEL

- Model versioning

- Model serving
- service
- Docker
- K8s

MONITORING
&
LOGGING

- Model decay trigger

MACHINE LEARNING PIPELINE

SOFTWARE CODE PIPELINE

CODE

- Trunk based dev.
- Code versioning

BUILD
&
INTEGRATION
TESTING

DEPLOY-
MENT
DEV
↓
PRODUCTION

FEEDBACK
new data from model performance

From Models to Decisions

A machine learning model is only useful when its output is **interpreted and acted upon**.

For example:

- A **fraud detection model** triggers a *review process*.
- A **credit risk model** informs *loan approval*.
- A **classification model** in healthcare supports *diagnosis decisions*.

Open points:

- Predictions are not final decisions — they are **decision-support tools**.
- Importance of **thresholds**, **confidence**, and **cost of false decisions**.
- Role of **human-in-the-loop** systems.

Visualizing Model Predictions

- Prediction Distribution
 - a. Show predicted probabilities (e.g., fraud probability, disease likelihood).
 - b. Use histograms, calibration plots, or scatter plots.
- Confusion Matrix and Thresholds
 - a. Demonstrate how different thresholds affect outcomes.
 - b. Visualization idea: a slider that adjusts decision threshold dynamically.
- Feature Importance
 - a. Visualize `feature_importances_` (tree models) or SHAP values.
 - b. Help users interpret why a prediction was made.

Interfaces for Machine Learning

Objective: Enable interaction between users and models.

Common interface types:

Type	Description	Example Tools
Web Apps	Interactive apps for model exploration	Streamlit, Gradio, Dash
Dashboards	Static or semi-interactive results	Power BI, Tableau
APIs	Backend-only services	Flask, FastAPI
Embedded Interfaces	Integration in other software	ERP systems, Chatbots

- Improves accessibility of ML results.
- Encourages collaboration between data scientists and stakeholders.
- Provides real-time feedback on model performance.

Introduction to Streamlit

What is Streamlit?

A Python-based open-source framework for creating interactive data apps and dashboards quickly.

Core Features:

- Build interfaces with pure Python (no HTML/JS).
- Supports interactive widgets (sliders, text boxes, buttons).
- Ideal for model demos, dashboards, and educational apps.

Installation:

```
pip install streamlit
```

Run an app:

```
streamlit run app.py
```



Live Demo: ML Decision Interface

Streamlit Example Code

```
import streamlit as st
import pickle
import numpy as np

# Load model
model = pickle.load(open('model.pkl', 'rb'))

st.title("Credit Card Fraud Detection App")

st.sidebar.header("Transaction Input Features")
amount = st.sidebar.number_input("Transaction Amount (€)", min_value=0.0)
v1 = st.sidebar.slider("Behavior Score (V1)", -5.0, 5.0, 0.0)
v2 = st.sidebar.slider("Risk Factor (V2)", -5.0, 5.0, 0.0)


input_data = np.array([[v1, v2, amount]])
prediction = model.predict(input_data)[0]
prob = model.predict_proba(input_data)[0][1]

st.subheader("Prediction Result")
if prediction == 1:
    st.error(f"⚠️ Fraud detected (probability: {prob:.2f})")
else:
    st.success(f"✅ Legitimate transaction (probability: {prob:.2f})")
```

Project:

Online Weather Prediction with Streamlit and River

Demonstrating Real-Time Machine Learning and Concept Drift Detection



Project: online Weather Prediction with Streamlit and River

- Traditional ML models are *trained offline* — they can't adapt quickly to new data.
- In real environments (weather, finance, IoT), **data arrives continuously**.
- Models must **learn incrementally** and detect **changes in patterns** (*concept drift*).

Goal:

Predict temperature in real time using **wind speed** data, continuously updating the model as new data arrives.

Project: online Weather Prediction with Streamlit and River

Component	Description
Data Source	Open-Meteo API provides live weather data (temperature, wind speed).
Model	Online Linear Regression from River (<code>LinearRegression</code> + <code>StandardScaler</code>).
Metrics	MAE (Mean Absolute Error) for performance tracking.
Drift Detection	ADWIN algorithm detects changes in error distribution.
Interface	Streamlit app for real-time visualization and control.

References:

- Bishop, C. M. (2006). Pattern Recognition and Machine Learning. Springer.
- Vapnik, V. N. (1995). The Nature of Statistical Learning Theory. Springer.
- Mitchell, T. M. (1997). Machine Learning. McGraw-Hill Education.
- Domingos, P. (2012). A Few Useful Things to Know About Machine Learning. Communications of the ACM, 55(10), 78–87.
- Jordan, M. I., & Mitchell, T. M. (2015). Machine Learning: Trends, Perspectives, and Prospects. Science, 349(6245), 255–260.